

Lesson Plan

Calico Installation

Summary

1. Introduction to Calico Installation Methods
2. Install Calico

Introduction to Calico Installation Methods

There are many different ways to install Calico. The Calico docs site includes recommended install options across a range of environments, so you don't need to be an expert to get started.

Broadly there are 4 different approaches to install.

Manifest

This is the most basic method for installing Calico. The Calico docs include a range of manifests for different environments. If you are an advanced user, you can customize the manifests to give you ultimate flexibility and control over your installation.

Operator

Calico 3.15 introduces the option to install Calico using an open-source operator, created by Tigera. This offers simplification for installing and configuring Calico without needing to customize manifests. Additionally, the operator allows you to have a uniform, self-healing environment. Using the Tigera operator is highly recommended.

Managed Kubernetes Services

Support for Calico is included with many of the most popular managed Kubernetes services (e.g. EKS, AKS, GKE, IKS), either enabled by default, or optionally enabled using the cloud provider's management consoles or command line tools, depending on the specific managed Kubernetes service.

Kubernetes Distros and Installers

Many Kubernetes distros and installers include support for installing Calico. (e.g. kops, kubespary, microk8s, etc). Most of these currently use manifest based installs under the covers.

Installing Calico for the Lab

We will be using the Tigera Operator to install and configure Calico.

If you are not already on host1, you can enter host1 by using the multipass shell command.

```
multipass shell host1
```

The command below will install the operator onto our lab kubernetes cluster:

```
kubectl create -f  
https://docs.projectcalico.org/archive/v3.16/manifests/tigera-operator.yaml
```

Validating the Operator installation

Following the operator being installed, we will validate that the operator is running:

```
kubectl get pods -n tigera-operator
```

The output from this command should indicate that the operator pod is running:

NAME	READY	STATUS	RESTARTS	AGE
tigera-operator-64f448dfb9-d2fdq	1/1	Running	0	2m33s

Installing Calico

After the operator is in a Running state, we will configure an Installation kind for Calico, specifying the IP Pool that we would like below.

Note that throughout the course we make use of inline manifests (piping stdin to kubectl) to make it easier for you to follow what each manifest does. In most cases it would be a more normal practice to use a vanilla kubectl command with a manifest file (e.g. kubectl apply -f my-installation.yaml). We recommend taking a minute to read through and make sure you understand the contents of each manifest we apply in this way throughout the rest of the course to get the most out of each example.

```
cat <<EOF | kubectl apply -f -
```

```

apiVersion: operator.tigera.io/v1
kind: Installation
metadata:
  name: default
spec:
  calicoNetwork:
    containerIPForwarding: Enabled
    ipPools:
      - cidr: 198.19.16.0/21
      natOutgoing: Enabled
      encapsulation: None
EOF

```

Validating the Calico installation

Following the configuration of the installation resource, Calico will begin deploying onto your cluster. This can be validated by running the following command:

```
kubectl get tigerastatus/calico
```

The output from the command when the installation is complete is:

NAME	AVAILABLE	PROGRESSING	DEGRADED	SINCE
calico	True	False	False	10m

We can review the environment now by invoking:

```
kubectl get pods -A
```

Example output:

NAMESPACE	NAME	READY	STATUS
tigera-operator	tigera-operator-84c5c5d6df-zb49b	1/1	Running
calico-system	calico-typha-868bb997ff-l22n7	1/1	Running
calico-system	calico-typha-868bb997ff-fvmws	1/1	Running
calico-system	calico-typha-868bb997ff-8qt45	1/1	Running
calico-system	calico-node-r94mp	1/1	Running

0	4m6s			
calico-system		calico-node-w5ptt	1/1	Running
0	4m6s			
calico-system		calico-node-zgrvb	1/1	Running
0	4m6s			
kube-system		helm-install-traefik-t68vd	0/1	
Completed	0	35m		
kube-system		metrics-server-7566d596c8-pccvz	1/1	Running
2	35m			
kube-system		local-path-provisioner-6d59f47c7-gh97b	1/1	Running
2	35m			
calico-system		calico-kube-controllers-89cf65556-c7gz7	1/1	Running
3	4m6s			
kube-system		coredns-7944c66d8d-f4q6g	1/1	Running
0	35m			
kube-system		svclb-traefik-9bxg2	2/2	Running
0	32s			
kube-system		svclb-traefik-pb72f	2/2	Running
0	32s			
kube-system		svclb-traefik-l6mzn	2/2	Running
0	32s			
kube-system		traefik-758cd5fc85-8hcdx	1/1	Running
0	32s			

Reviewing Calico pods

Let's take a look at the Calico pods that have been installed by the operator.

```
kubectl get pods -n calico-system
```

Example Output:

NAME	READY	STATUS	RESTARTS	AGE
calico-typha-5d788c654b-56wp9 4h28m	1/1	Running	0	
calico-node-2bkv5 4h28m	1/1	Running	0	
calico-kube-controllers-5dcfdb5f4-vpgx5 4h28m	1/1	Running	0	
calico-node-8465h 4h26m	1/1	Running	0	
calico-typha-5d788c654b-wn7gf 4h24m	1/1	Running	0	
calico-node-qmq5j 3h57m	1/1	Running	0	
calico-typha-5d788c654b-rd8kl 3h56m	1/1	Running	0	

From here we can see that there are different pods that are deployed:

- **calico-node:** Calico-node runs on every Kubernetes cluster node as a DaemonSet. It is responsible for enforcing network policy, setting up routes on the nodes, plus managing any virtual interfaces for IPIP, VXLAN, or WireGuard.
- **calico-typha:** Typha is as a stateful proxy for the Kubernetes API server. It's used by every calico-node pod to query and watch Kubernetes resources without putting excessive load on the Kubernetes API server. The Tigera Operator automatically scales the number of Typha instances as the cluster size grows.
- **calico-kube-controllers:** Runs a variety of Calico specific controllers that automate synchronization of resources. For example, when a Kubernetes node is deleted, it tidies up any IP addresses or other Calico resources associated with the node.

Reviewing Node Health

Finally, we can review the health of our Kubernetes nodes by invoking the `kubectl` command.

```
kubectl get nodes -A
```

Example output:

NAME	STATUS	ROLES	AGE	VERSION
control	Ready	master	37m	v1.18.10+k3s1
node2	Ready	<none>	16m	v1.18.10+k3s1
node1	Ready	<none>	31m	v1.18.10+k3s1

Now we can see that our Kubernetes nodes have a status of Ready and are operational. Calico is now installed on your cluster and you may proceed to the next module: Installing the Sample Application.